

Evaluation: Hashi-Solver (Group M)

Noel Huibers, Laura Schröder, Florian Würmseer

22.11.2023

1 Introduction

In the context of Hashi puzzles, also known as Bridges, we aim to evaluate the scalability of our encoding method as grid sizes increase. This assessment will measure the performance of our solver and explore ways to make the encoding more efficient. Our objective is to ensure that our solver can efficiently handle larger grids while maintaining accuracy. The following sections will present our analysis of the solver's performance using Criterion and the steps we've taken to optimize its efficiency.

2 Dataset

The benchmarks and measurements were designed to evaluate the solver's performance across a range of puzzles with varying grid sizes and complexities. The following test files were employed, each characterized by a specific grid size and layout of the puzzle:

- **test1.txt** - Featuring a grid size of 5x5. It includes a straightforward puzzle layout with minimal islands and bridges, ideal for assessing the solver's efficiency in handling simpler, less demanding puzzles.
- **test8.txt** - This file presents an intermediate level of complexity with a 9x9 grid. It introduces a more challenging puzzle layout, including a greater number of islands and more intricate bridge connections.
- **test20.txt** - Featuring a grid size of 15x15. A puzzle layout with a dense arrangement of islands and bridges. It tests the solver's capability to efficiently process and solve large-scale, intricate puzzles that demand significant computational power and advanced problem-solving strategies.

Each test file progressively increases in complexity and grid size, providing a comprehensive spectrum for evaluating the solver's performance. This diverse dataset aims to examine the solver's scalability, efficiency, and robustness in handling puzzles ranging from simple to highly complex. The goal is to identify strengths, limitations, and potential areas for optimization in the solver's design and implementation.

3 Performance Measurement: Solver Modes

A comparative analysis of the `encode_mode`, `solve_mode`, and `esr_mode` functions was conducted. The benchmark was performed multiple times to ensure reliability of the results.

Test Case	Average Execution Time (ns)	Performance Change
encode/test1	448.07	+0.26% to +0.50%
encode/test8	454.19	-0.36% to +1.72%
encode/test20	451.30	-0.43% to +1.16%
solve/test1	458.95	-3.69% to -0.85%
solve/test8	461.68	-1.41% to +0.02%
solve/test20	462.67	-0.85% to +0.47%
esr/test1	462.34	-1.69% to +0.08%
esr/test8	458.21	-0.61% to +0.51%
esr/test20	450.55	-0.16% to +0.96%

Table 1: Benchmark Results

The benchmark results provide insights into the performance of various backend functions under different test conditions. The following observations can be made:

- **Consistency in Execution Time:** The average execution times for the `encode` and `esr` functions across different test files are relatively consistent, indicating stable performance across different input sizes.
- **Performance Changes:** Notable performance changes are observed in the `solve` function, especially for `test1`, where the execution time shows a significant decrease. This might suggest optimizations or efficiencies in the algorithm that become more apparent with certain types of input.
- **Impact of Test File Size:** Larger test files (e.g., `test20`) do not show a significant increase in execution time, which is indicative of efficient algorithm scalability with larger input sizes.
- **Outliers:** The presence of outliers in several test cases, particularly in `test8` and `test20` for the `esr` function, could indicate potential inconsistencies in performance under certain conditions. These outliers might warrant further investigation to understand their cause.
- **Overall Stability:** The majority of the performance changes are within a small margin, suggesting overall stability in the function implementations. However, the variance in the `encode` function for `test8` and `esr` function for `test8` and `test20` could be areas to explore for potential optimizations.

In conclusion, the benchmark results show a generally stable and efficient performance across the tested functions, with some areas of potential optimization and a need for further investigation into the causes of outliers.

4 Performance Measurement: Helper-Functions

Function	Test1 (μ s)	Test8 (μ s)	Test20 (μ s)
parse_input	862.78	615.68	1843.81
generate	268.67	875.07	1596.31
generate_dimacs	113.71	148.44	241.44
solve	16.26	12.22	15.49
write_solution	97.07	76.26	75.93
reconstruct_puzzle	81.50	168.97	323.24

Table 2: Average Execution Times of Functions

These results, obtained through multiple executions, provide valuable insights into the performance characteristics of different functions within our solver. The key observations from the analysis are as follows:

- The `parse_input` function shows a notable increase in execution time with the increasing complexity and size of the test cases, indicating its sensitivity to the input size.
- In contrast, the `solve` function maintains relatively consistent and low execution times across all test cases, suggesting its efficiency is less affected by the complexity of the puzzle.
- The `generate` and `generate_dimacs` functions exhibit a moderate increase in execution times with larger puzzles, which aligns with their roles in puzzle generation and DIMACS format conversion.
- The `write_solution` and `reconstruct_puzzle` functions demonstrate a balanced performance across different test cases, although a slight increase in time is noticeable in more complex puzzles for `reconstruct_puzzle`.

These findings highlight the areas where optimization efforts could be focused, particularly in the `parse_input` and `generate` functions for larger puzzles. Additionally, the overall efficiency and scalability of the `solve` function under varying conditions are noteworthy. In conclusion, the benchmarking exercise has successfully identified performance trends and potential bottlenecks, setting the stage for targeted improvements in the solver’s performance.

5 Optimization Opportunities

Based on the analyzed results, several potential areas for performance optimization have been identified. Recommendations for enhancing the efficiency of various functions are provided below.

5.1 Using u8 instead of usize

Switching from `usize` to `u8` could theoretically reduce memory usage, as `u8` occupies 1 byte compared to `usize`'s 8 bytes on 64-bit systems. This might lead to improved performance with large data sets by allowing more data to fit into the CPU cache.

Important considerations:

- `u8` can only store values between 0 and 255. If values for `x`, `y`, or `connections` exceed 255, `u8` is unsuitable.
- CPUs are most efficient with values matching their native word size, which is `usize` or `u64` on a 64-bit system. Processing `u8` values might be slower due to the extra work required by the CPU.
- Memory usage is often not the primary performance constraint. The switch to `u8` may not lead to a significant performance improvement if the bottleneck is elsewhere in the code.

The performance changes observed after using `u8` instead of `usize` are summarized in the following table:

Test Case	Average Execution Time (ns)	Performance Change
encode/test1	456.16	-4.10% to -1.22%
encode/test8	456.77	-1.83% to -0.32%
encode/test20	463.06	-3.30% to -1.21%
solve/test1	470.19	-0.31% to +1.16%
solve/test8	466.53	-2.05% to -0.35%
solve/test20	458.68	-0.77% to +0.19%
esr/test1	457.76	-0.84% to -0.06%
esr/test8	454.86	-1.33% to -0.12%
esr/test20	448.92	-0.77% to -0.39%

Table 3: Performance Changes with `u8` instead of `usize`

The benchmark results indicate a mixed impact on performance due to the use of `u8` instead of `usize`. The `encode` test cases generally show an improvement in performance, with a decrease in execution times across all test sizes. This suggests that the memory efficiency gained from using `u8` positively affects these operations. However, the `solve` and `esr` test cases demonstrate a less consistent

trend. Some tests indicate slight improvements, while others show marginal regressions or no significant change. This variation suggests that while memory efficiency is enhanced, other factors, possibly related to CPU processing efficiency or algorithmic complexity, are influencing the overall performance. Overall, the use of `u8` shows potential benefits in terms of memory usage and execution time for certain operations, but its impact varies depending on the specific functions and operations being performed.

5.2 Further potentially beneficial changes for performance improvement

In this section, additional possibilities for improving the performance of the Hashi-Solver are suggested:

- **Using BufRead instead of read_to_string:** BufRead allows reading the file line-by-line, potentially reducing memory usage and increasing speed when processing small parts of a file.
- **Avoiding unnecessary clone operations:** In Rust, cloning is expensive as it creates a full copy of the data object. Minimizing or eliminating clone operations, especially in lines like `bridge in bridges.clone()`, could enhance performance.
- **Using more efficient data structures:** Replacing Vec with HashMap or HashSet can be advantageous for frequent element searches, as these structures provide faster search operations.

6 Conclusion

The benchmark results reveal a generally stable and efficient performance of the solver, with certain areas identified for potential optimization. The comparative analysis of different solver modes – `encode_mode`, `solve_mode`, and `esr_mode` – indicates consistent execution times and highlights the scalability and robustness of the algorithm, especially notable in the `solve` function's performance across different puzzle complexities.

The investigation into optimization opportunities, such as using `u8` instead of `usize` and implementing more efficient data structures or input processing methods, shows mixed impacts. While some changes result in improved execution times and memory efficiency, others have less consistent effects, suggesting that the solver's performance is influenced by a combination of factors, including memory usage, CPU processing efficiency, and algorithmic complexity.